

The Representation Stack

A Reference Model for How Reality Becomes Measurable, Represented, Optimized, and Interpreted

A. Jacobs | Reality Drift Framework

Introduction

Systems that operate at scale cannot act on reality in its full continuity and context. They act on selected signals, measurements, indicators, models, reports, and explanations. The Representation Stack describes the transformations by which an underlying condition becomes something a system can measure, compare, represent, optimize, and interpret.

Each transformation is necessary. Each also changes what passes through it. Context is removed, ambiguity is resolved, proxies acquire authority, and representations become easier to operate on than the realities they were built to represent.

The Stack

Reality provides the ground condition. Measurement selects observable signals. Metrics aggregate those signals into comparable proxies. Operational representation organizes those proxies into models, dashboards, classifications, targets, and other forms that can guide action. Optimization directs behavior toward the objectives encoded in those representations. Narrative constructs meaning around the resulting outputs.

What scales is not reality itself. What scales is the representation.

The Core Distinction

The stack is not a claim that measurement, metrics, models, or narratives are inherently false. Representation is unavoidable. The question is whether each layer remains open to correction from the reality beneath it.

A representation can become more coherent, legible, and persuasive while becoming less responsive to the condition it was built to represent. Coherence is not contact.

Four Properties of the Stack

Selection Is Not Neutral - Measurement begins by deciding what can enter the system. Unmeasured conditions do not disappear, but they lose operational visibility. What the system cannot register becomes difficult to evaluate, reward, or correct.

Accuracy Is Local - A value may accurately describe its selected signal while failing to preserve the broader meaning or condition that signal stands for. Accuracy at one layer does not guarantee fidelity across the stack.

Optimization Changes the Object - Once an operational representation guides action, people and systems adapt to the objectives encoded within it. The representation begins shaping the reality it was created to describe. Its relationship to the underlying condition can therefore weaken even while represented performance improves.

Narrative Closes Uncertainty - Narrative turns partial outputs into a stable account of what happened, why it happened, and what should happen next. This enables coordination, but it can also conceal the transformation history beneath the explanation.

The Layers

Reality is the reference condition beneath the stack. The five layers track how selected aspects of that condition become measurable, comparable, operable, governable, and interpretable. Real systems move between these layers recursively. Metrics alter measurement, narratives change what institutions collect, and optimized outputs become new inputs.

Ground Condition: Reality

Function: The condition the system is intended to perceive, serve, explain, or change.

What it contains: Full causal and contextual relations, including conditions that have not been formalized.

System limitation: Reality cannot be captured in full. It remains the reference against which the alignment of every subsequent layer must be assessed.

Layer 1: Measurement

Function: Selects observable signals and records them in a stable form.

What it preserves: Chosen features, observations, events, counts, or traces.

What it loses: Unselected context, tacit knowledge, causal texture, timing, and conditions the instrument cannot register.

Substitution: The measurable signal stands in for the larger condition.

Characteristic failure: Orientation failure. The system sees what its instruments admit and becomes insensitive to what they exclude.

Layer 2: Metrics

Function: Aggregates measurements into indicators that support comparison, evaluation, and coordination.

What it preserves: Regularity, rank, direction, rate, threshold, or another standardized relation.

What it loses: Granularity, edge cases, qualitative difference, uncertainty, and relations outside the aggregation rule.

Substitution: A comparable proxy stands in for a complex state or outcome.

Characteristic failure: Proxy failure. The metric remains valid as a calculation while weakening as a description of the thing that matters.

Layer 3: Operational Representation

Function: Organizes measurements and metrics into dashboards, models, interfaces, classifications, targets, or simulations that can guide action.

What it preserves: Selected relationships, system state, and a usable view of complexity.

What it loses: Process visibility, provenance, ambiguity, hidden assumptions, and the distinction between observed and inferred content.

Substitution: The operable model stands in for the condition from which it was derived.

Characteristic failure: Symbolic inversion. The representation becomes the object the system responds to, even when contrary ground signals remain available.

Layer 4: Optimization

Function: Uses metrics and operational representations to guide decisions, incentives, allocation, and behavior.

What it preserves: The objective encoded by the target and the performance differences visible through the representation.

What it loses: Original intent, unmodeled tradeoffs, externalities, and forms of value that cannot improve the represented target.

Substitution: Improving the representation stands in for improving the underlying condition.

Characteristic failure: Target capture. Participants and systems adapt to the represented objective, weakening its relationship to reality.

Layer 5: Narrative

Function: Constructs explanation, causal sequence, identity, legitimacy, and shared understanding around optimized representations and their outputs.

What it preserves: A communicable account that coordinates interpretation and action.

What it loses: Contradiction, uncertainty, alternate causality, unresolved complexity, and the history of transformation.

Substitution: A coherent account stands in for a contested and partially observed condition.

Characteristic failure: Closure failure. The explanation becomes more stable than the evidence and filters later observations through itself.

What Each Transition Gains and Loses

Each transformation solves a coordination problem. The same move creates a characteristic form of distance. Drift is best diagnosed by examining what each layer makes possible and what the system can no longer see or correct.

Reality to Measurement

Primary gain: Visibility.

Typical loss: Context and unselected signals.

Drift risk: What cannot be recorded becomes operationally absent.

Measurement to Metrics

Primary gain: Comparability.

Typical loss: Granularity and qualitative difference.

Drift risk: The proxy inherits authority from the observations beneath it.

Metrics to Operational Representation

Primary gain: Operability.

Typical loss: Process, provenance, and uncertainty.

Drift risk: The representation appears complete because its exclusions and construction history are no longer visible.

Operational Representation to Optimization

Primary gain: Direction.

Typical loss: Original purpose and unmodeled tradeoffs.

Drift risk: The system improves represented performance faster than the underlying condition.

Optimization to Narrative

Primary gain: Meaning and coordination.

Typical loss: Ambiguity, competing explanations, and the history of transformation.

Drift risk: A coherent account of optimized outputs is mistaken for evidence of real-world improvement.

The Central Substitution

Reality becomes signal. Signal becomes proxy. Proxy becomes an operational representation. The representation becomes the object of optimization. Optimized outputs become the basis of explanation. At every step, the new object is easier to move through the system than the condition it replaces.

The risk is not that the new object is unreal. It is real as a measurement, metric, model, target, optimized output, or narrative. The risk is category collapse. The system forgets what kind of object it is operating on and grants the representation the authority of the reality beneath it.

How Representation Becomes Drift

The stack describes the architecture through which reality becomes measurable, represented, optimized, and interpreted. The sequence below identifies the mechanisms that can turn these ordinary transformations into self-reinforcing misalignment. These are not additional layers. They describe how drift develops across the stack.

- 1. Filtering:** Some signals become visible while others disappear from operational view.
- 2. Compression:** Selected signals are reduced into portable metrics, categories, or proxies.
- 3. Representation Substitution:** The operational representation becomes the object the system responds to.
- 4. Proxy Optimization:** Behavior shifts toward the targets and objectives encoded within that representation.
- 5. Constraint Weakening:** Conditions outside the representation lose the power to redirect the system.

6. Coherence Stabilization: Outputs increasingly agree with one another and appear complete.

7. Feedback Weakening: Contradictory evidence arrives late, indirectly, or without enough authority to change the model.

8. Drift Amplification: Representations are reused as inputs, allowing inherited assumptions and errors to compound.

The decisive change occurs when representation stops functioning as a revisable interface with reality and becomes the environment within which the system evaluates itself.

Once constraint weakening and feedback weakening occur, later layers can validate earlier ones. Narratives justify targets, dashboards confirm narratives, and optimized outputs become training data or institutional precedent. Recursive reuse gives inherited assumptions the appearance of independent confirmation.

Trace the Decision Backward

Apply the stack to one consequential output such as a score, recommendation, dashboard state, policy, model response, or institutional conclusion. Begin at the point of action and move toward the ground condition.

Step 1: Name the Operational Object

Identify the representation the system is actually responding to.

Step 2: Recover the Transformation History

Identify the measurements, aggregations, targets, model choices, and interpretive steps that produced it.

Step 3: Identify Excluded Conditions

Ask what can matter in reality without appearing in the representation.

Step 4: Locate the Active Substitution

Find where a signal, metric, model, or explanation has become equivalent to the thing it represents.

Step 5: Test Corrective Contact

Identify an observation that could force revision and determine whether it can reach the decision point.

Step 6: Reintroduce Grounding

Restore direct observation, qualitative judgment, provenance, uncertainty, counter-metrics, or outcome feedback at the layer where correction was lost.

A system has not been grounded merely because it contains human review. Grounding requires a feedback path capable of changing the representation, target, or process that produced the decision.

The Architecture Beneath Survivable Wrongness

Reality Drift occurs when a system remains operational and internally coherent while gradually losing corrective contact with the conditions it represents or serves. The Representation Stack explains how this can happen without a single dramatic failure.

The absence of failure is not proof of alignment. A system can continue producing scores, reports, decisions, and explanations because each layer remains internally functional. Drift becomes visible when those outputs are compared with the conditions they were meant to track, preserve, or improve.

Model Boundary

The Representation Stack does not replace domain knowledge, causal analysis, or empirical validation. It identifies where those forms of correction must enter and asks whether a system remains responsive to what its representations stand for.

Within the wider framework, Semantic Fidelity evaluates whether meaning survives transformation. Reality Drift names the accumulated misalignment that develops when corrective contact no longer keeps pace with those transformations.

Representations are necessary, but alignment depends on preserving the ability to return through them.

Representation enables scale. Feedback preserves contact.