

AI Alignment: The Problem Is Not Behavior, It Is Meaning

Reality Drift Failure Modes Series | 01 Semantic Decoupling

A. Jacobs — Reality Drift Framework

Core Claim

AI alignment is not primarily about behavior. It is about semantic fidelity. The challenge is preserving the meaning of human intent as that intent is translated into representations, objectives, and outputs.

The Pattern

Most people treat AI alignment as a behavioral problem centered on bias, safety, and guardrails. But the deeper issue is semantic. Systems fail because they execute instructions without truly capturing the intent behind them. Alignment is ultimately about keeping meaning anchored to reality as it moves through models, metrics, and optimization.

At first, this distinction seems subtle. If a system produces acceptable outputs, it appears aligned. If it avoids harm and follows instructions, it appears to be working.

But behavior is only the surface. What matters is whether the system remains connected to the conditions that give those instructions meaning.

The Mechanism

Every instruction given to an AI system is a compressed representation of intent.

That intent must pass through multiple layers:

- Language (how the instruction is expressed)
- Model representation (how it is interpreted internally)
- Optimization (how the system decides what to produce)
- Output (what is generated)

At each step, meaning can shift. Systems do not “understand” intent in a grounded sense. They process patterns. As a result, they can execute instructions in ways that are internally consistent but externally misaligned.

This creates a structural problem. The system can appear aligned at the level of behavior while drifting at the level of meaning.

Where This Shows Up

This shows up whenever a system has to act on a representation of intent rather than intent itself. As information moves through layers of translation, compression, and optimization, preserving meaning becomes increasingly difficult.

A language model may produce a fluent and helpful-looking response while missing what the user was actually trying to ask. The same pattern appears in automated systems that optimize metrics standing in for goals, or reinforcement systems that maximize reward while drifting from the purpose the reward was meant to approximate.

In each case, the system follows the instruction as represented, not as intended.

Reality Drift Connection

Reality Drift describes what happens when systems continue operating while losing alignment with reality.

AI alignment failures are one expression of this broader pattern. As instructions move through layers of representation, they are compressed, interpreted, and optimized in ways that can detach them from their original grounding. The system remains functional. But the meaning has shifted.

A system is not aligned because it follows an instruction. It is aligned only when the meaning behind the instruction survives.

Core Framework and Sources

- [Substack \(Articles\)](#)
- [GitHub \(Full Library\)](#)
- [DOI \(Research Paper\)](#)
- [Glossary & Definition](#)