

When Rational Behavior Produces Unreal Outcomes

Incentive Misalignment, Principal-Agent Problems, and Reality Drift

Reality Drift Framework — Bridge Brief #2

A. Jacobs

The Established Insight

Every organization depends on delegation.

Owners rely on managers.

Citizens rely on governments.

Patients rely on healthcare systems.

Shareholders rely on executives.

Employers rely on employees.

The challenge is straightforward.

How can people pursuing their own decisions reliably produce someone else's intended outcome?

Economics and organizational theory developed one of the most influential answers through the **principal-agent problem**.

Its central insight is simple:

The goals of the agent may differ from the goals of the principal.

When incentives diverge, behavior changes.

The agent begins optimizing for different objectives than the ones the system was originally created to achieve.

The result is not necessarily irrational behavior.

It is rational behavior directed toward a different objective.

What Incentive Theory Explains

Incentive theory explains why people optimize.

Behavior responds to rewards, penalties, career advancement, status, compensation, and institutional pressure.

When these incentives change, behavior changes with them.

This framework explains:

- principal-agent problems
- moral hazard
- agency costs
- regulatory capture
- bureaucratic self-interest
- rent seeking
- governance failures

The sequence is familiar:

delegated purpose → incentive structure → rational behavior → unintended outcome

Nothing requires the participants to be dishonest.

They simply respond to the environment created by the system.

This remains one of the strongest explanations for organizational behavior ever developed.

The Boundary

Many modern systems drift even when no obvious incentive failure exists.

Participants may genuinely want to do good work.

The incentives may be internally coherent.

No individual is acting maliciously.

Yet the organization gradually becomes less responsive to the reality it was created to serve.

Performance indicators improve.

Reporting becomes more sophisticated.

Optimization becomes more effective.

The institution remains operational.

Its purpose becomes harder to recognize.

This raises a different question.

Why can systems drift even when rational people are responding appropriately to rational incentives?

The source paper argues that incentives alone do not explain these failures because organizations rarely optimize reality directly.

They optimize representations of reality.

The Reality Drift Extension

Reality Drift shifts attention from incentives to the objects incentives operate upon.

Organizations rarely interact directly with the realities they exist to serve.

Instead they coordinate through:

- reports
- dashboards
- forecasts
- metrics
- models
- performance indicators
- credentials

These representations make large organizations possible.

They also become the objects around which incentives are organized.

The sequence becomes:

reality → representation → incentive → optimization → organizational behavior

At first, the representation faithfully reflects reality.

Over time, incentives increasingly reward improvement in the representation itself.

The report becomes easier to improve than the underlying condition.

The dashboard becomes easier to optimize than the experience it summarizes.

The forecast becomes easier to defend than the uncertainty it once represented.

The incentive has not failed.

Its target has changed.

Representational Incentives

A **representational incentive** is an incentive attached primarily to a symbolic representation rather than to the underlying condition it was created to describe.

This distinction matters.

An employee may honestly pursue a KPI.

A researcher may honestly optimize a benchmark.

A manager may honestly improve reporting accuracy.

An institution may honestly satisfy every procedural requirement.

The representation improves.

Whether reality improved becomes increasingly difficult to determine.

The organization has not stopped optimizing.

It has become progressively more responsive to the representation through which reality is being observed.

The Bridge

Incentive theory and Reality Drift explain different levels of the same system.

Incentive theory asks:

Why do people optimize particular objectives?

Reality Drift asks:

What happens when the objectives themselves become increasingly representational?

Incentive theory explains behavior.

Reality Drift examines the fidelity of what behavior is organized around.

This creates a distinct possibility.

Perfectly rational actors can collectively produce increasingly unrealistic systems.

No one intends the drift.

Each participant responds appropriately to the symbolic environment they inherit.

The representations become progressively more influential than the realities they were meant to coordinate.

Where the Distinction Matters

Bureaucracies

Employees satisfy procedural requirements.

Documentation improves.

The institution becomes increasingly effective at producing compliant representations of work.

Whether the public is better served becomes less visible.

Healthcare

Clinicians complete documentation, coding, and reporting requirements accurately.

Administrative performance improves.

The patient's lived experience becomes increasingly indirect.

Artificial Intelligence

Engineers optimize benchmark performance and reward functions.

The optimization succeeds.

The relationship between evaluation scores and real-world capability becomes harder to evaluate.

Higher Education

Universities expand credentials, rankings, reporting systems, and measurable outputs.

Every participant behaves rationally.

The relationship between credentials and capability gradually weakens.

Core Distinction

Incentive theory explains why people optimize the objectives placed before them.

Reality Drift examines what happens when those objectives become increasingly organized around representations rather than the realities they were intended to serve.

Systems rarely drift because people stop acting rationally. They drift because rational behavior becomes progressively coordinated around representations whose relationship to reality is quietly weakening.